

# AGEC 652 - Lecture 2.1

## Numerical arithmetic

Diego S. Cardoso

Spring 2022

# Course roadmap

1. Intro to Scientific Computing
2. **Numerical operations and representations**
  1. **Numerical arithmetic** ← You are here
  2. Numerical differentiation and integration
  3. Function approximation
3. Systems of equations
4. Optimization
5. Structural estimation

\*These slides are based on Bryant & O'Hallaron (2011), Miranda & Fackler (2001), Judd (1998), and course materials by Ivan Rudik.

# Briefest review of Computer Systems

(or How do computers work?)

# Why bother?

Knowing the basics of how computers perform the instructions we give them is essential to

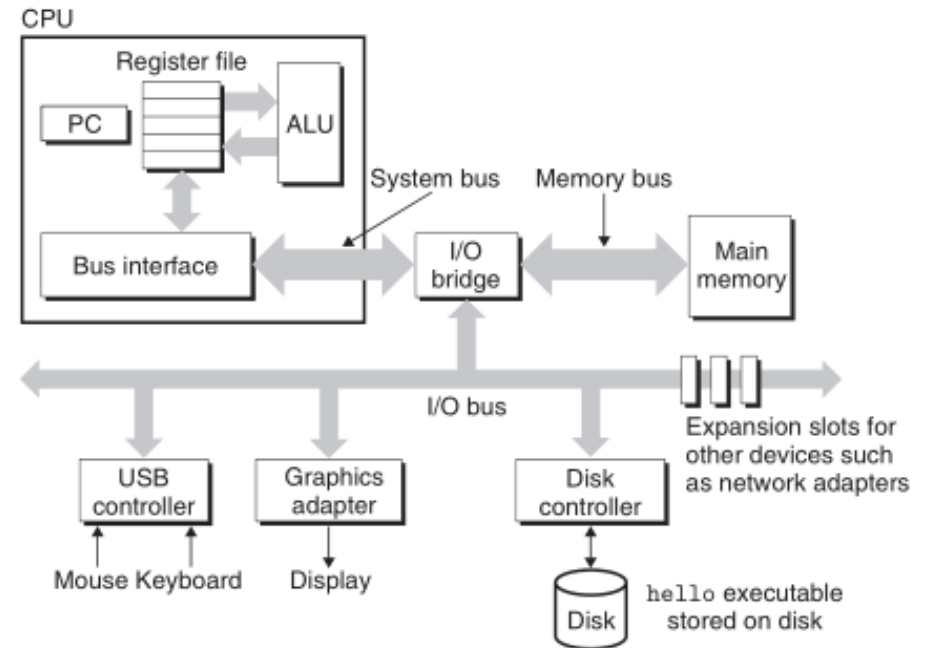
- Understand why we write code in certain ways
- Why your program fails or takes forever to run
- Write more efficient routines
- Understand the limits of computational methods in quantitative research

# The building blocks of a computer

**Buses:** where information flows between components. They are design to transport fixed-size chunks known as *words* (32 or 64 bytes in most computers)

**I/O devices:** connect the system to the external world

- Input: keyboard, mouse, disks, networks
- Output: display, disks, networks




# The building blocks of a computer

**Main memory:** temporary storage of *instructions* (i.e. programs) *and data*

- Physically, a collection of DRAM chips
- Logically, a **linear array of bytes** (8 bits), each with its *unique address*

When we allocate a variable of type `Int64`, our program will determine an available address (where that variable starts in memory) and reserve a size (64 bits = 8 bytes, in this case)

When we assign `x = 10`, our program writes the binary equivalent of integer `10` in that memory address



# The building blocks of a computer

```
x = [1 2 3]
```

```
## 1×3 Matrix{Int64}:  
##  1  2  3
```

```
# This is how we get the physical address (0x means it's in hexadecimal base)  
pointer_from_objref(x)
```

```
## Ptr{Nothing} @0x000001fb2a8be6d0
```

16 digits of hexadecimal number (4 bits):  $16 \times 4 = 64$ -bit address

```
# Get the size  
sizeof(x)
```

```
## 24
```

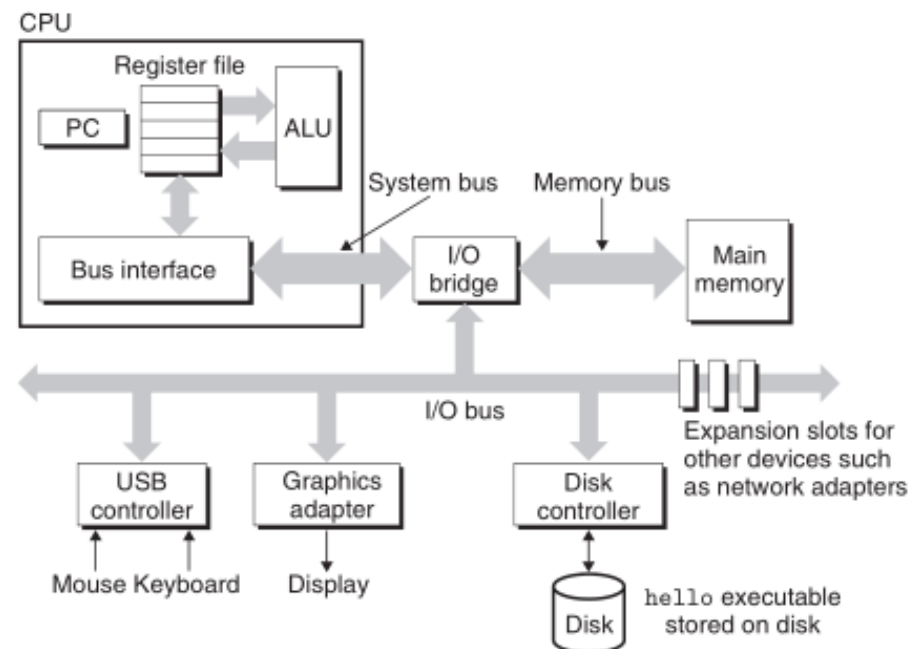
$3 \times \text{Int64}$  (8 bytes) = 24 bytes



# The building blocks of a computer

**CPU**, or processor: it's the computer's engine. It has 3 parts:

- *Program Counter* (**PC**): keeps the memory address of the current instruction
- **Registers**: tiny but *ultra-fast memory* (e.g.: 16 64-bit registers on Intel i7)
- *Arithmetic/logic unit* (**ALU**): performs operations over the values stored in registers



# The building blocks of a computer

CPUs generally operate with a limited number of simple instructions. Any complex calculation translated by the compiler into many simple steps

There are four main types of instructions

- *Load*: copy a value from main memory into a register
- *Store*: copy a value from a register to main memory
- *Operate*: copy 2 values from register to ALU, perform an arithmetic operation on them, store the result in a register
- *Jump*: point the PC to an instruction in another memory address

# What happens when we run a program?

```
function hello()  
  print("hello, world\n")  
end;  
  
hello()
```

```
## hello, world
```

# What happens when we run a program?

1) Our function `hello` is compiled: another program will read our human language and convert it to *machine language* (instructions)

- I.e., our code goes from a sequence of bits (0s and 1s) that can be represented as characters in a display to another sequence of bits that represents instructions which the **CPU** can understand

2) The compiled function is placed into the **main memory**

- Compiled code can also be stored in permanent memory (hard drives, SSD, disks). Those are the *executable* files. But before they are executed, they must be loaded into the main memory anyway

# What happens when we run a program?

3) We type `hello()` in Julia REPL. Our **input** prompts Julia and the Operating System (OS) to prepare those instructions for execution

- E.g.: by *jumping* the PC to the address where they are stored

4) The **CPU** executes those instructions and **outputs** the results back in the command line, which is then converted into image by a graphics adapter and our display

**Moral of the story: the computer spends A LOT of time just moving information around and very little time actually doing "real work!"**

# Running fast

The computer spends A LOT of time just moving information around and very little time actually doing "real work!" What to do?

Efficient computation needs to

1. Minimize the need to pass information around
  - This is what we do, for example, when we preallocate variables, avoid temporary allocations, vectorize
2. Make information travel faster
  - Faster hardware (not always affordable/available)
  - **Cache memory!**

# Cache memory

Physics law: bigger memory = slower memory

- Your hard disk can be  $1,000\times$  larger than the main memory but it might take  $10,000,000\times$  longer for the processor to read

But the main memory is too big/slow. The CPU actually has its own internal layers of fast memory with different sizes

- Modern processors have registers + 3 levels of cache memory

# Cache memory

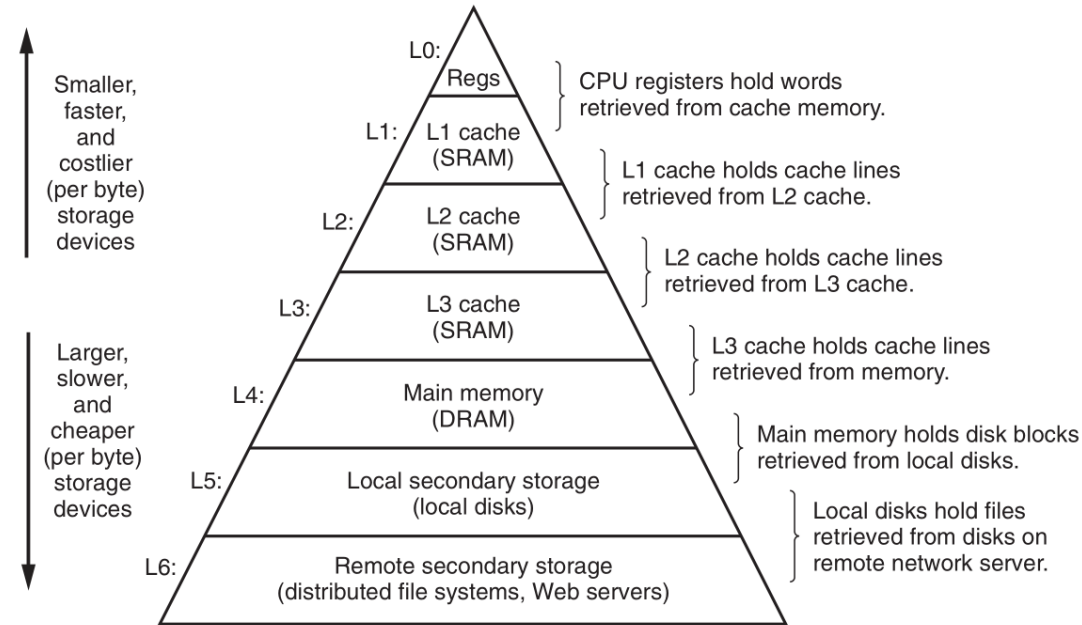


Table 2-4. Cache Parameters of the Skylake Microarchitecture

Level	Capacity / Associativity	Line Size (bytes)	Fastest Latency <sup>1</sup>	Peak Bandwidth (bytes/cyc)	Sustained Bandwidth (bytes/cyc)	Update Policy
First Level Data	32 KB/ 8	64	4 cycle	96 (2x32B Load + 1*32B Store)	~81	Writeback
Instruction	32 KB/8	64	N/A	N/A	N/A	N/A
Second Level	256KB/4	64	12 cycle	64	~29	Writeback
Third Level (Shared L3)	Up to 2MB per core/Up to 16 ways	64	<sup>44</sup>	32	~18	Writeback

*These are Intel specifications →*



# Writing cache-friendly code


Now that you knowing about the memory hierarchy, you can use that knowledge to write *more efficient, cache-friendly code*. That means:

- Making the common case go fast
  - Programs often spend most of the time in a few loops. **Focus on the inner loops of your core functions**
- Minimize the chances of needing to bring variables from slow memory
  - Reference local variables
  - If you need to reference the same variable in multiple iterations, think about a way order your loops in a way that you make those **repeated references in sequence**

# Operating System

There's an important component between your software and the hardware: the *Operating System* (OS)

The OS is a software layer that manages and makes accessible all the hardware components. It transforms those components in simple abstractions that can be easily manipulated with much simpler software instructions



# Numerical Arithmetic

# Why bother?

Simple arithmetic:  $(10^{-20} + 1) - 1 = 10^{-20} + (1 - 1)$ , right?

Let's check:

```
x = (1e-20 + 1) - 1;  
y = (1e-20) + (1 - 1);  
x == y
```

```
## false
```

## What happened?!

# Why bother?

Welcome to the world of **finite precision**

```
x
```

```
## 0.0
```

```
y
```

```
## 1.0e-20
```

Numbers are an ideal/abstract concept that works perfectly in analytic operations

Computers are a physical machine with limited resources: they can't have infinite precision

Understanding this limitation is crucial in numerical analysis

# Number storage

**Q: How are numbers physically represented by a computer?**

*A: In a binary (or Base 2) system. This means each digit can only take on 0 or 1*

The system we're most used to is Base 10: each digit can take on 0-9

**Q: So which numbers can be represented by a computer?**

*A: A **subset** of the rational numbers*

But computers have **finite** memory and hard disk space, there are **infinite** rational numbers

This imposes a strict limitation on the storage of numbers

# Number storage

Numbers are stored as:  $\pm mb^{\pm n}$

- $m$  is the *significand*
- $b$  is the *base*,
- $n$  is the *exponent*

All three are integers

- The *significand* typically gives the significant digits
- The *exponent* scales the number up or down in magnitude

# Number storage

The size of numbers a computer can represent is limited by how much space is allocated for a number

Space allocations are usually 64 bits: 53 for  $m$  and 11 for  $n$

We've seen these types before

```
println(typeof(5.0))
```

```
## Float64
```

```
println(typeof(5))
```

```
## Int64
```



# Number storage

- `Int64` means it is a `integer` with 64 bits of storage
- `Float64` means it is a `floating point number` with 64 bits of storage
  - Floating point just means  $b^{\pm n}$  can move the decimal point around in the significand

These two types take the same space in memory (64 bits) but are interpreted differently by the processor

# The limits of computers

Limitations on storage suggest three facts

**Fact 1:** There exists a **machine epsilon**: the smallest relative quantity representable by a computer

**Machine epsilon** is the smallest  $\epsilon$  such that a machine can always distinguish

$$N + \epsilon > N > N - \epsilon$$

# The limits of computers

```
println("Machine epsilon  $\epsilon$  is $(eps(Float64))")
```

```
## Machine epsilon  $\epsilon$  is 2.220446049250313e-16
```

```
println("Is  $1 + \epsilon/2 > 1$ ? $(1.0 + eps(Float64)/2 > 1.0)")
```

```
## Is  $1 + \epsilon/2 > 1$ ? false
```

The function `eps` gives the distance between 1.0 and the next representable `Float64`

# The limits of computers

```
println("The smallest representable number larger than 1.0 is $(nextfloat(1.0))")
```

```
## The smallest representable number larger than 1.0 is 1.00000000000000002
```

```
println("The largest representable number smaller than 1.0 is $(prevfloat(1.0))")
```

```
## The largest representable number smaller than 1.0 is 0.9999999999999999
```

# The limits of computers

The machine epsilon changes depending on the amount of storage allocated

```
println("32 bit machine epsilon is $(eps(Float32))")
```

```
## 32 bit machine epsilon is 1.1920929e-7
```

```
println("Is  $1 + \epsilon/2 > 1$ ? $(Float32(1) + eps(Float32)/2 > 1)")
```

```
## Is  $1 + \epsilon/2 > 1$ ? false
```

**There is a trade-off between precision and storage requirements**

- This matters for low-memory systems like GPUs

# The limits of computers

## Fact 2: There is a **smallest representable number**

```
println("64 bit smallest float is $(floatmin(Float64))")
```

```
## 64 bit smallest float is 2.2250738585072014e-308
```

```
println("32 bit smallest float is $(floatmin(Float32))")
```

```
## 32 bit smallest float is 1.1754944e-38
```

```
println("16 bit smallest float is $(floatmin(Float16))")
```

```
## 16 bit smallest float is 6.104e-5
```

# The limits of computers

## Fact 3: There is a **largest representable number**

```
println("64 bit largest float is $(floatmax(Float64))")
```

```
## 64 bit largest float is 1.7976931348623157e308
```

```
println("32 bit largest float is $(floatmax(Float32))")
```

```
## 32 bit largest float is 3.4028235e38
```

```
println("16 bit largest float is $(floatmax(Float16))")
```

```
## 16 bit largest float is 6.55e4
```

# The limits of computers: time is a flat circle

```
println("The largest 64 bit integer is ${typemax(Int64)}")
```

```
## The largest 64 bit integer is 9223372036854775807
```

```
println("Add one to it and we get: ${typemax(Int64)+1}")
```

```
## Add one to it and we get: -9223372036854775808
```

```
println("It loops us around the number line: ${typemin(Int64)}")
```

```
## It loops us around the number line: -9223372036854775808
```



# The limits of computers

## The scale of your problem matters

If a parameter or variable is  $> \text{floatmax}$  or  $< \text{floatmin}$ , you will have a very bad time

Scale numbers appropriately (e.g. millions of dollars, not millionths of cents)

# Error

We can only represent a finite number of numbers

This means we will have error in our computations

Error comes in two major forms:

1. Rounding
2. Truncation

# Rounding

We will always need to round numbers to the nearest computer representable number. This introduces error

```
println("Half of  $\pi$  is:  $\pi/2$ ")
```

```
## Half of  $\pi$  is: 1.5707963267948966
```

The computer gave us a rational number, but  $\pi/2$  should be irrational

# Truncation

Lots of important numbers are defined by infinite sums  $e^x = \sum_{n=0}^{\infty} \frac{x^n}{n!}$

It turns out that computers can't add up infinitely many terms because there is finite space  $\rightarrow$  we need to truncate the sum

# Why does this matter?

**Errors are small, who cares?**

**You should!**

Because errors can propagate and grow as you keep applying an algorithm (e.g. function iteration)

# Error example 1

Let's get back to the example we saw earlier

```
x = (1e-20 + 1) - 1
```

```
## 0.0
```

```
y = 1e-20 + (1 - 1)
```

```
## 1.0e-20
```

```
println("The difference is: $(x-y).")
```

```
## The difference is: -1.0e-20.
```

# Error example 1

Why did we get  $y > x$ ?

- For  $(10^{-20} + 1) - 1$ : when we added  $10^{-20}$  to 1, it got rounded away
- For  $10^{-20} + (1 - 1)$ : here  $1 - 1$  was evaluated first and return 0, as we would expect

**Adding numbers of greatly different magnitudes does not always works like you would want**

# Error example 2

Consider a simple quadratic eq.  $x^2 - 26x + 1 = 0$  with solution  $x = 13 - \sqrt{168}$

```
println("64 bit: 13 - √168 = $(13-sqrt(168))")
```

```
## 64 bit: 13 - √168 = 0.03851860318427924
```

```
println("32 bit: 13 - √168 = $(convert(Float32,13-sqrt(168)))")
```

```
## 32 bit: 13 - √168 = 0.038518604
```

```
println("16 bit: 13 - √168 = $(convert(Float16,13-sqrt(168)))")
```

```
## 16 bit: 13 - √168 = 0.0385
```



# Error example 2

Let's check whether they solve the equation

```
x64 = 13-sqrt(168); x32 = convert(Float32,13-sqrt(168)); x16 = convert(Float16,13-sqrt(168));  
f(x) = x^2 - 26x + 1;  
println("64 bit: $(f(x64))")
```

```
## 64 bit: 7.66053886991358e-15
```

```
println("32 bit: $(f(x32))")
```

```
## 32 bit: 0.0
```

```
println("16 bit: $(f(x16))")
```

```
## 16 bit: 0.0004883
```

# Error example 3

Let's just subtract two numbers:  $100000.2 - 100000.1$

We know the answer is  $0.1$

```
println("100000.2 - 100000.1 is: $(100000.2 - 100000.1)")
```

```
## 100000.2 - 100000.1 is: 0.099999999999126885
```

```
if (100000.2 - 100000.1) == 0.1
  println("and it is equal to 0.1")
else
  println("and it is not equal to 0.1")
end
```

```
## and it is not equal to 0.1
```

# Error example 3

Why do we get this error?

Neither of the two numbers can be precisely represented by the machine!

$$100000.1 \approx 8589935450993459 \times 2^{-33} = 100000.0999999999767169356346130$$

$$100000.2 \approx 8589936309986918 \times 2^{-33} = 100000.1999999999534338712692261$$

So their difference won't necessarily be 0.1

## Error example 3

There are tools for approximate equality. Remember the `$\approx$` operator we saw last lecture? You can use it

```
(100000.2 - 100000.1) ≈ 0.1 # You type \approx then hit TAB
```

```
## true
```

This is equivalent to:

```
isapprox(100000.2 - 100000.1, 0.1)
```

```
## true
```

**This matters, particularly when you're trying to evaluate logical expressions of equality**

# Course roadmap

This concludes the first lecture of Unit 2. Up next:

1. Intro to Scientific Computing
2. **Numerical operations and representations**
  1. Numerical arithmetic
  2. **Numerical differentiation and integration** ←
  3. Function approximation
3. Systems of equations
4. Optimization
5. Structural estimation